

NoSuggest

Data Protection Impact Assessment

Prepared under GDPR Article 35 guidance, self-assessed, not independently reviewed

Version 1.1 | 7 October 2026 | Author: Varun Manjunath, Founder of NoSuggest

A note on this document

I wrote this myself, as the sole person who builds and operates [NoSuggest](#). It isn't reviewed by a lawyer or a data protection professional, and I say that plainly rather than let the document imply otherwise. Its purpose is to reason through what NoSuggest does with data, in enough detail that a school, a researcher, or an advocacy group can see the reasoning rather than just my conclusion. Wherever I'm relying on a claim I haven't independently verified, I've marked it as such.

1. What this assessment covers

This DPIA covers [nosuggest.com](#) and the [NoSuggest](#) web app: the frontend interface, the private backend worker that proxies' requests to the YouTube Data API, and Cloudflare Web Analytics. It does not cover YouTube's own embedded video player, YouTube's own data practices, or Cloudflare's data practices as a network provider more broadly. Those are third-party services outside my control, and I address them as dependencies and risks rather than as something I can assess on their behalf. It also covers Cloudflare Turnstile, which the backend now uses to verify that API requests come from real users rather than automated scripts. Turnstile's own data handling is outside my control and is addressed as a dependency below.

2. Description of the processing

Data	Where it goes	Why it's processed
Channel ID, channel handle, or search keyword the user types	Sent from the browser to my private backend worker, which forwards the query to the YouTube Data API and returns the result	This is how the app finds and displays channels and videos. It's the core function of the product.
IP address	Passes through Cloudflare's network as part of any web request. Cloudflare's own documentation states it is used transiently for country-level geolocation within Web Analytics, and is not persisted in the analytics report itself	Not something I request or store myself. I'm relying on Cloudflare's stated handling.
Aggregate page-view and performance data	Collected by Cloudflare Web Analytics, a cookie-free beacon script	Understanding basic traffic and site performance. Retention is Cloudflare's platform default of one month; I haven't changed it and there's no per-site setting to change it.
Kids Mode PIN	Stored only in the browser's local storage, in plain text. Never sent to any server.	Lets a parent or teacher restrict channel changes on a shared device. This is local device state, not data I receive or process.
Followed channels and saved videos	Stored only in the browser's local storage	Lets the app remember a user's choices without an account. Not transmitted to or stored by me.
Cloudflare Turnstile verification signals (IP	Exchanged directly between the user's browser and Cloudflare's Turnstile service when the app	Bot protection: prevents automated scripts from exhausting the YouTube Data API

Data	Where it goes	Why it's processed
address, browser characteristics)	contacts the backend API. NoSuggest does not receive or store these signals. The backend receives only a short-lived, cryptographically signed session token (for feed and cheap lookups) or a single-use verification token (for video search and channel name search), both issued by the backend itself after Turnstile confirms the request is human. Both token types contain no personal information and expire within five minutes or on first use.	quota that NoSuggest relies on to function for all users. The signals are processed by Cloudflare, not by me, and are subject to Cloudflare's Turnstile Privacy Addendum.

3. Is this personal data, and who is the controller

There is no account, no persistent user profile, and no data linking a query to a returning visitor.

On who the controller is: schools and other institutions that use NoSuggest are given no configuration access, no admin panel, and no ability to change what analytics or backend I use. That makes me the data controller for NoSuggest's own processing, and a school or institution using it is best described as a deployer of the service, not a joint controller.

4. Necessity and proportionality

- The channel ID, handle, or search query is the minimum needed to perform the function the user asked for. There's no reasonable way to look up a channel or search for a video without sending the search term somewhere.
- No account and no persistent profile are created. This is a deliberate design choice, not an omission, since the app doesn't need to remember a user across devices or sessions to function.
- Aggregate analytics are used only to understand traffic and performance, not to build any individual profile. I could operate without analytics entirely; I've chosen to keep a minimal, privacy-oriented version rather than none, to understand per country, per device, per referrer traffic.
- Cloudflare Turnstile is used to distinguish real users from automated scripts. The signals it collects (IP address, browser characteristics) are the minimum Cloudflare requires to make that determination. NoSuggest receives none of those signals directly. The backend uses a cost-proportionate model: expensive API calls (video search, channel name search, each costing 100 times more YouTube Data API quota units) require a fresh single-use Turnstile token per request, while cheaper calls (feed loads, channel lookups by ID or handle, each costing 0–1 units) reuse a short-lived session token valid for five minutes. This limits unnecessary Turnstile challenges for routine use while protecting the highest-cost operations.

5. Risk identification and mitigation

Risk	Likelihood and severity	Mitigation
A search query typed by a child could, in rare cases, contain sensitive personal content (for example, a name or a health-related term used as a search).	Low likelihood, and I don't store search queries beyond the momentary request needed to return results.	The backend architecture ensures complete separation between user queries and IP data: search parameters are processed statelessly, and the backend proxy is configured to ignore incoming IP headers, preventing any association between search terms and client IP addresses.

Risk	Likelihood and severity	Mitigation
Cloudflare's stated handling of IP addresses in Web Analytics is a third-party claim I haven't independently verified.	Low likelihood of harm, but I can't currently confirm it beyond Cloudflare's own documentation.	I'm relying on Cloudflare's published privacy documentation. I'd revisit this if Cloudflare's own policy changes.
Cloudflare Turnstile processes browser signals (IP address, browser characteristics) to determine whether a request is human. This is a third-party data flow I don't control.	Low likelihood of harm. Turnstile is designed to be privacy-preserving and does not use tracking cookies. The signals are transient and used only for the bot-detection decision.	Disclosed in the privacy policy and in section 5 of the whitepaper. I rely on Cloudflare's Turnstile Privacy Addendum for their stated handling. NoSuggest receives no Turnstile signals directly; only a derived pass/fail decision is used to gate API access.
The embedded YouTube player and YouTube's own API are outside NoSuggest's control, and YouTube's own data practices apply once a video plays.	Outside my ability to mitigate directly.	I disclose this dependency clearly in the whitepaper and to any institution evaluating NoSuggest, rather than implying I control the full viewing experience.
Kids Mode is a deterrent aimed at preventing a young child from casually leaving a curated feed. It is not a security boundary. Anyone who clears the browser's local storage or opens the page in a private or incognito window, exits Kids Mode without needing the PIN at all. This applies wherever Kids Mode is described, including the whitepaper and any institutional materials, so it should be described accurately as convenience control rather than access control.		

6. Conclusion

I believe NoSuggest processes a narrow, low-sensitivity set of data, with no account, no persistent profile, and no data linking a query to an individual. None of this point to a high-risk processing activity that would require a full, formal DPIA under Article 35 rather than this screening-level assessment, but I'm stating that as my own read rather than a settled legal conclusion, since I'm not a lawyer or a data protection professional.

I intend to revisit this document if the architecture changes, if Cloudflare's analytics policy changes, or if a specific institution's deployment introduces new data flows I haven't accounted for here.